

A dark blue vertical bar runs down the left side of the page. A blue arrow points to the right from this bar, containing the date.

07/04/2019

# Assignment 2

Automated Cloud Services

Several thin, curved lines in dark blue and light grey originate from the bottom left and sweep upwards and to the right.

Jonathan McDonagh  
20074520

## Contents

<b>Introduction</b> .....	3
<b>Core Assignment Specification</b> .....	3
<b>Step 1 (Master Instance)</b> .....	3
<b>Step 2 (AMI)</b> .....	4
<b>Step 3 (VPC)</b> .....	4
<b>Step 4 (Launch Configuration)</b> .....	5
<b>Step 5 (Load Balancer)</b> .....	5
<b>Step 6 (Auto-Scaling Group)</b> .....	7
<b>Step 7 (Auto-Scaling Policy)</b> .....	9
<b>Step 8 (CloudWatch)</b> .....	10
<b>Step 9 (Test Traffic to the Load Balancer)</b> .....	12
<b>Step 10 (To Show the Load is Distributed Across the Web Servers)</b> .....	13
<b>Step 11 (Custom CloudWatch Metric)</b> .....	14
<b>References</b> .....	16

## Figures

Figure 1: Cloud Architecture Diagram .....	3
Figure 2: AMI.....	4
Figure 3: VPC.....	4
Figure 4: Launch Configuration.....	5
Figure 5: Load Balancer.....	5
Figure 6: Load Balancer (a) .....	6
Figure 7: Load Balancer (b) .....	6
Figure 8: Auto Scaling Group .....	7
Figure 9: cpu100.sh.....	8
Figure 10: ASG top .....	8
Figure 11: Auto Scaled Instances .....	9
Figure 12: Auto Scaling Policy (cpu50).....	9
Figure 13: Cloud-Watch Policies .....	10
Figure 14: cpu100 script for Cloud Watch .....	10
Figure 15: Cloud Watch Auto Scaling.....	11
Figure 16: Cloud Watch Instances Terminating .....	11
Figure 17: Cloud Watch Back to Min .....	12
Figure 18: curl command .....	12
Figure 19: Load Distribution.....	13
Figure 20: AWS Metric Script .....	14
Figure 21: Cloud Watch Monitoring (1) .....	14
Figure 22: Cloud Watch Monitoring Increased (2).....	15

## Introduction

The objective of this assignment is to demonstrate the deployment and automated management of a load-balanced auto-scaling web application. As you can see below here is my Cloud Architecture Diagram.

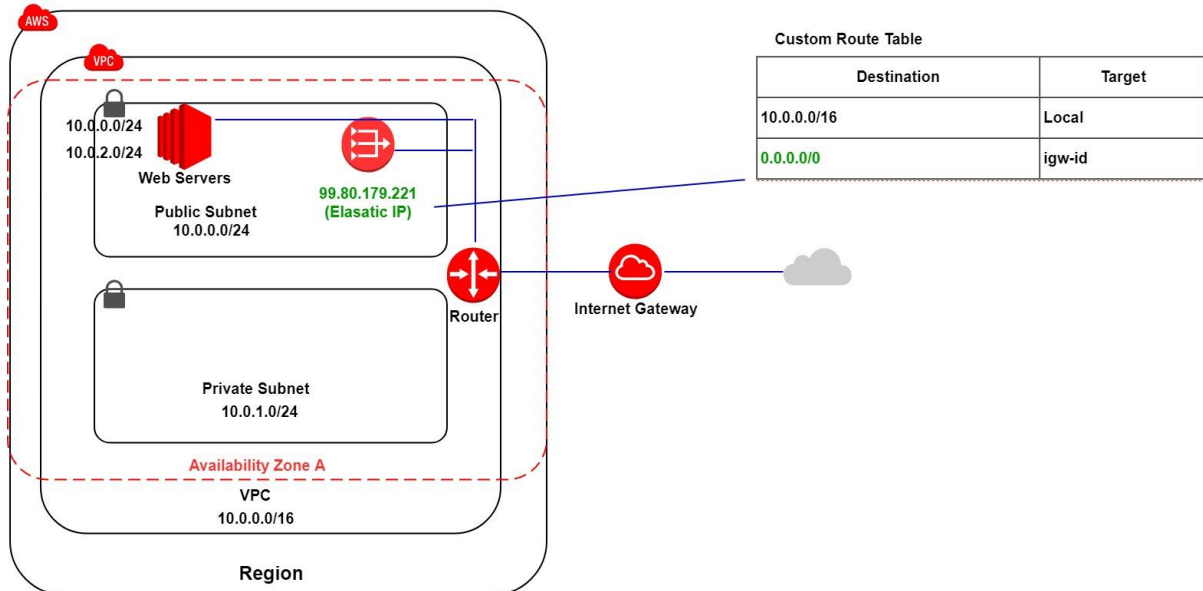


Figure 1: Cloud Architecture Diagram

## Core Assignment Specification

### Step 1 (Master Instance)

The creation and configuration of a “master” instance. I created an ec2 instance on AWS based on my VPC, in the user data I install Apache web server.

## ASSIGNMENT 2

### Step 2 (AMI)

The creation of the AMI, I clicked on the “master” instance I then clicked on the actions button and I went down to image and clicked Create, this created the AMI for the master instance.

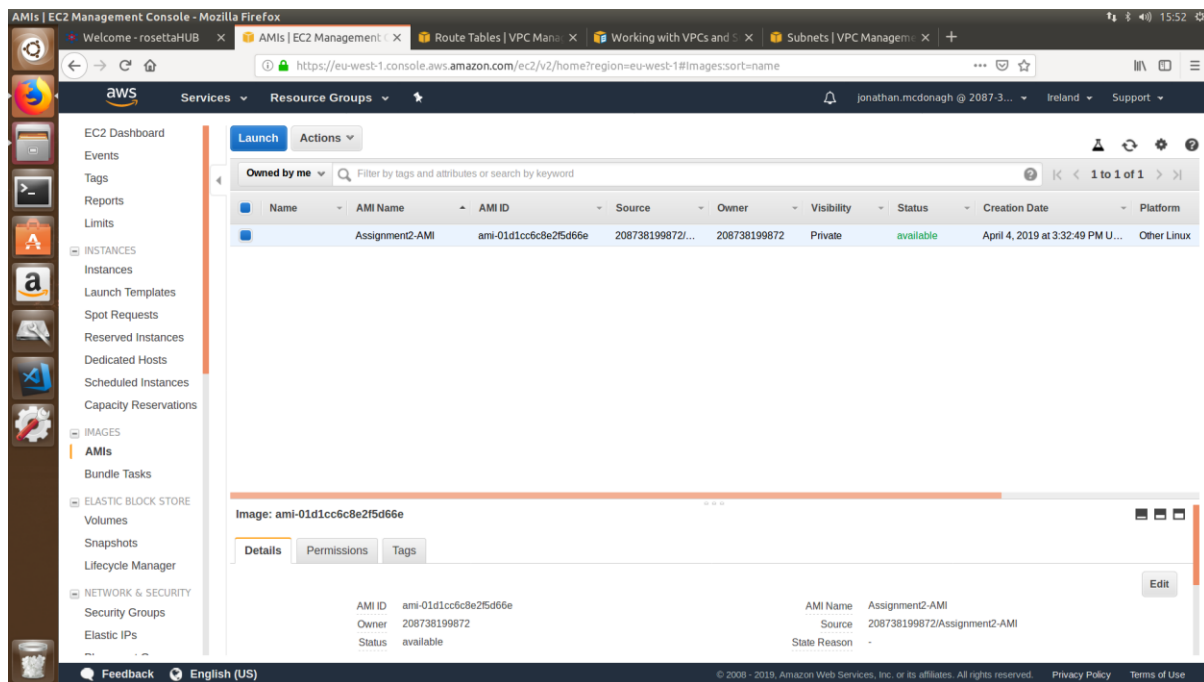


Figure 2: AMI

### Step 3 (VPC)

The creation of the VPC with public and private subnets. I created a VPC with a Public and Private Subnet. But before creating the VPC I had to allocate an Elastic IP for the VPC. After I created the VPC I created another Public Subnet. I had one Public Subnet in EU-West-1a and another Public Subnet in EU-West-1b. The name of my VPC is VPC-Assignment2.

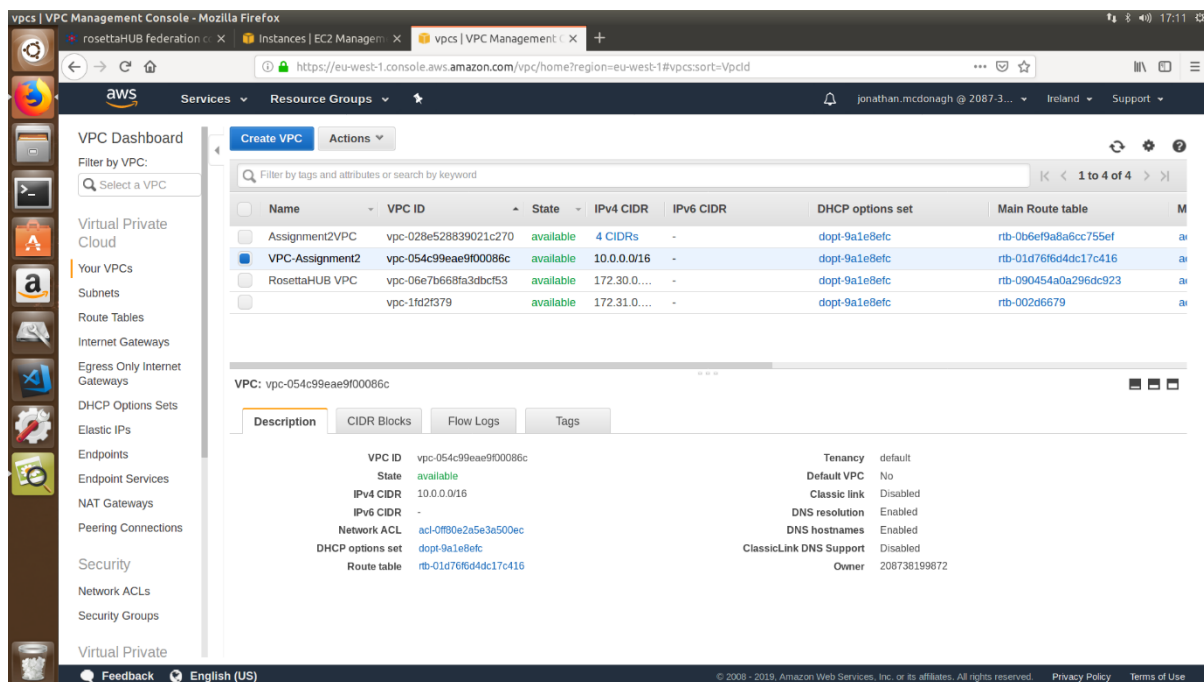


Figure 3: VPC

## Step 4 (Launch Configuration)

The creation of a launch configuration based on my AMI.

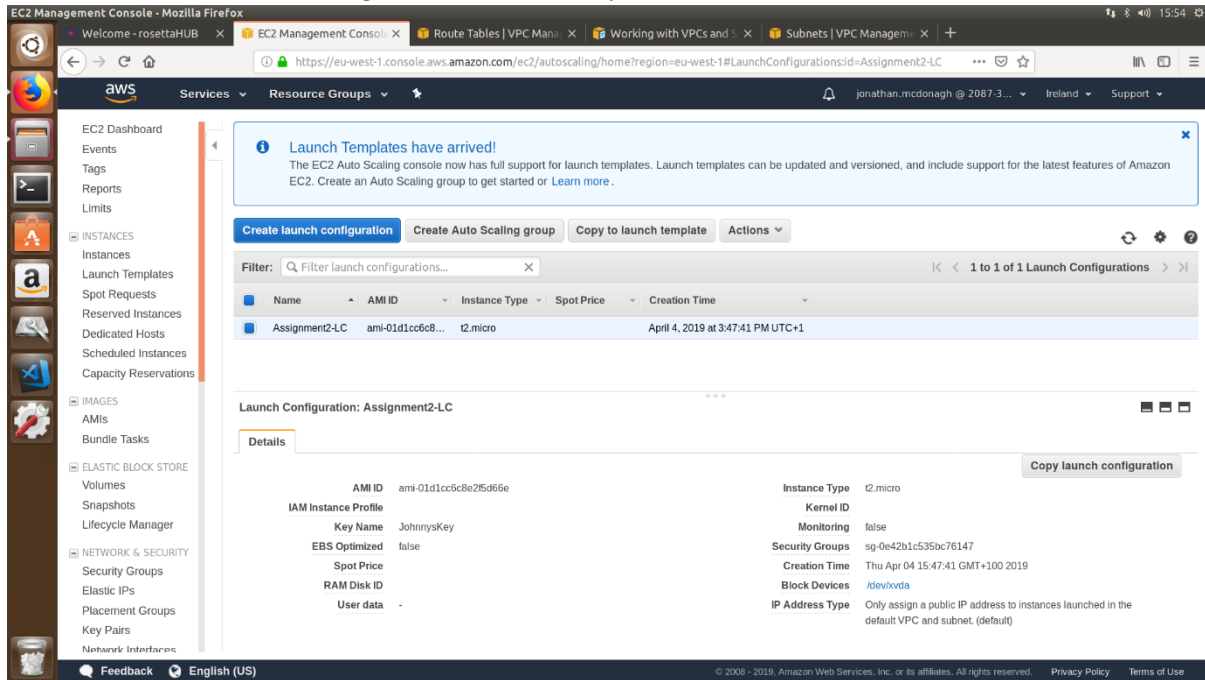


Figure 4: Launch Configuration

## Step 5 (Load Balancer)

The Creation of a load balancer. When I copied the DNS name onto the browser and refreshed it, the web server would change between a and b. The two different instances depending on the availability zone so A would show for EU-West-1a and B would show for EU-West-1b. As you can see below.

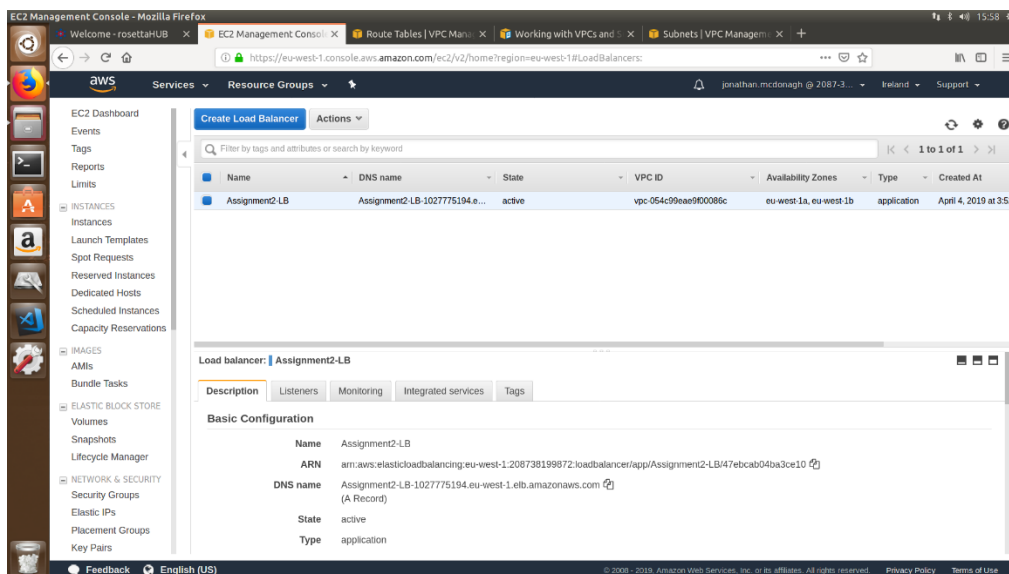


Figure 5: Load Balancer

## ASSIGNMENT 2

a

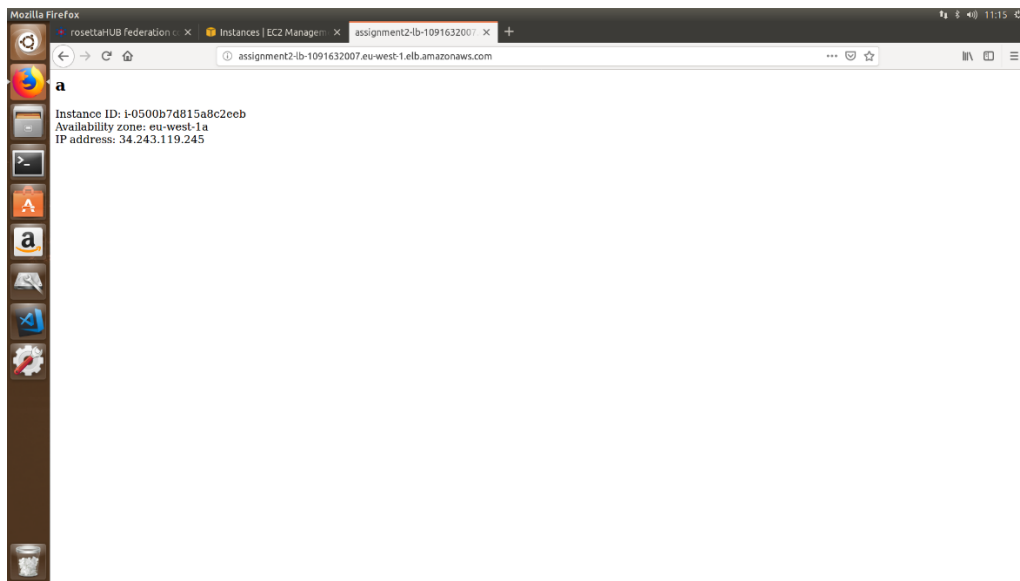


Figure 6: Load Balancer (a)

b

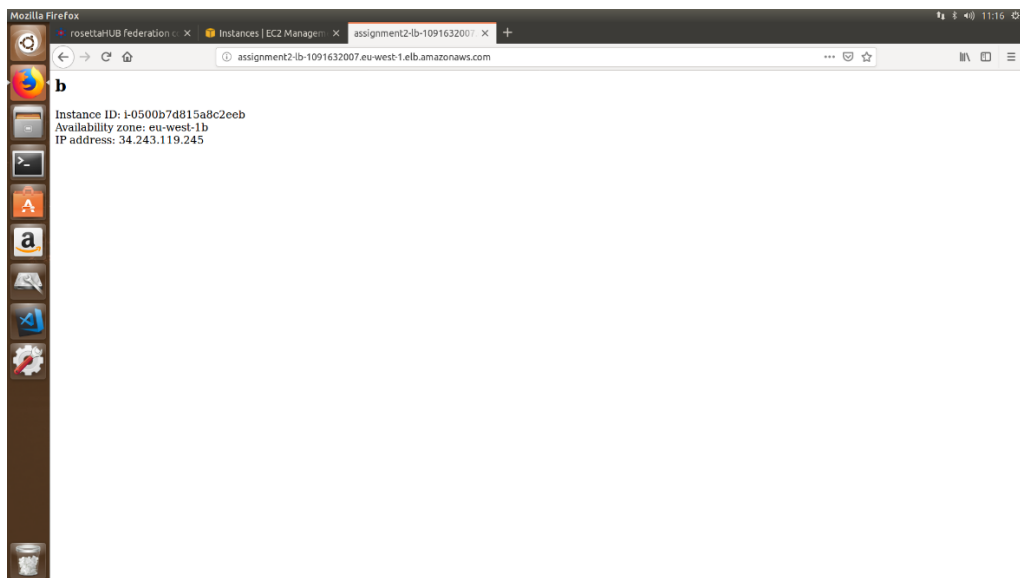


Figure 7: Load Balancer (b)

## Step 6 (Auto-Scaling Group)

To create the Auto-Scaling Group I first had to create the Launch Configuration based on my AMI. After this I created a Auto-Scaling Group based on my Launch Configuration. When creating my Auto-Scaling Group I choose my VPC. I then chose my two Public Subnets.

I then selected “Use scaling policies to adjust the capacity of this group” and set it up to scale between 2 and 4 instances when the Target Value was above 50. I gave the Instances 90 seconds to warm up after scaling. I also set the Target Group to my Target Group I created for the Load Balancer.

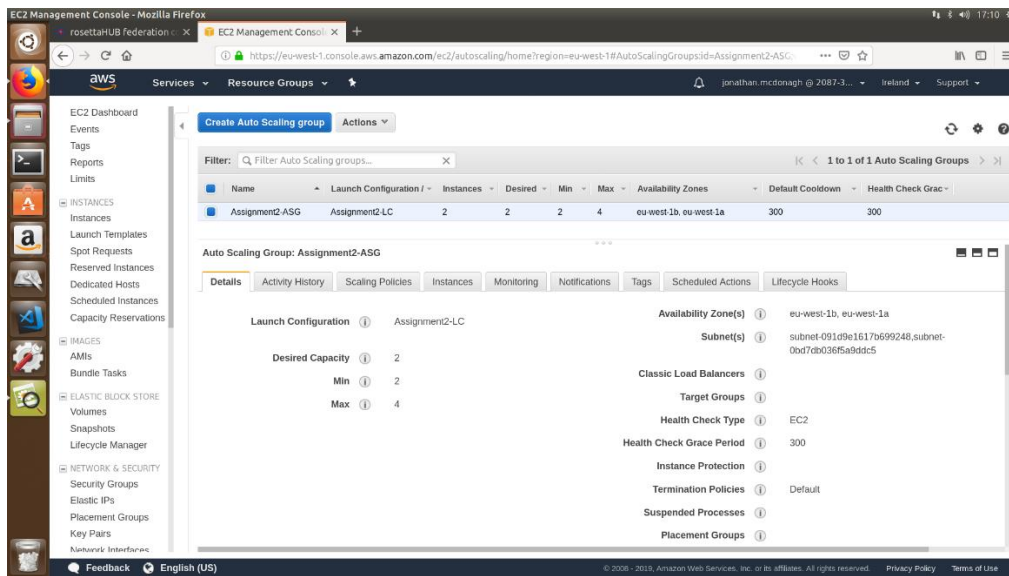


Figure 8: Auto Scaling Group



## ASSIGNMENT 2

To trigger an increase in group size I logged into my instance that was created from the Auto-Scaling Group and created a simple bash script called "cpu100.sh".

Inside the script it had the following code:

```
While true; do x=0; done
```

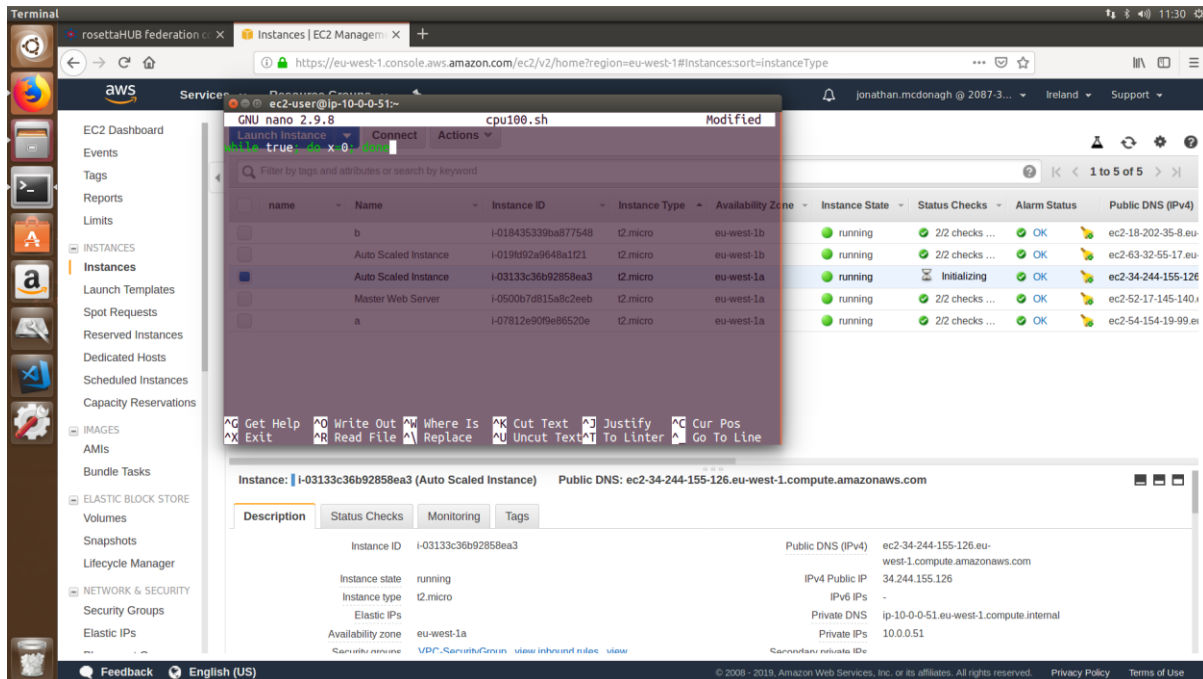


Figure 9: cpu100.sh

I then used the **top** command to view the process CPU utilisation. After a few minutes another instance started up.

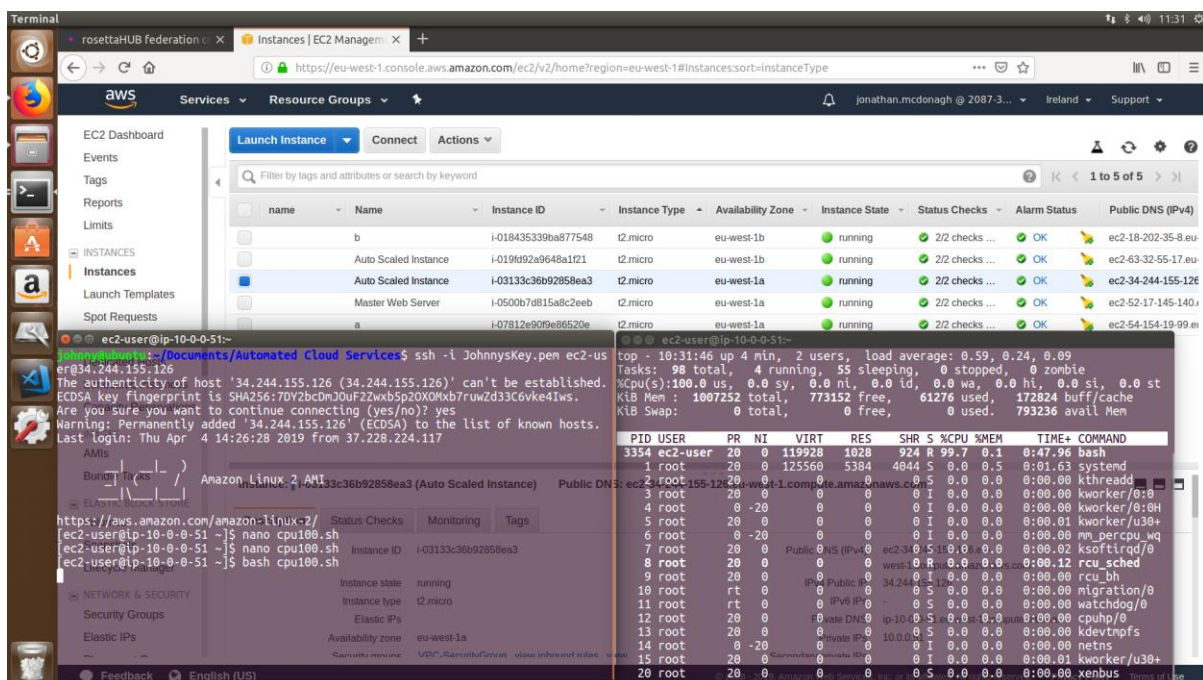


Figure 10: ASG top

## ASSIGNMENT 2

As you can see below two instances were automatically created in the auto scaling group.

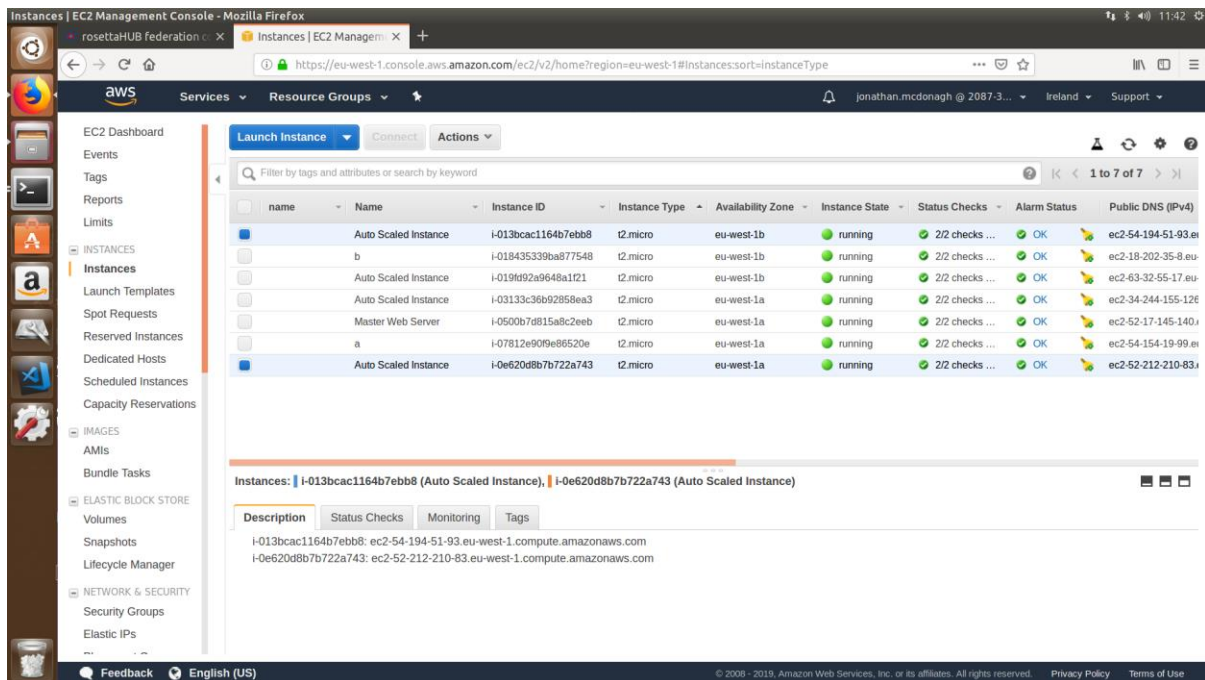


Figure 11: Auto Scaled Instances

### Step 7 (Auto-Scaling Policy)

As you can see below here is my Auto Scaling Policy for my Auto Scaling Group.

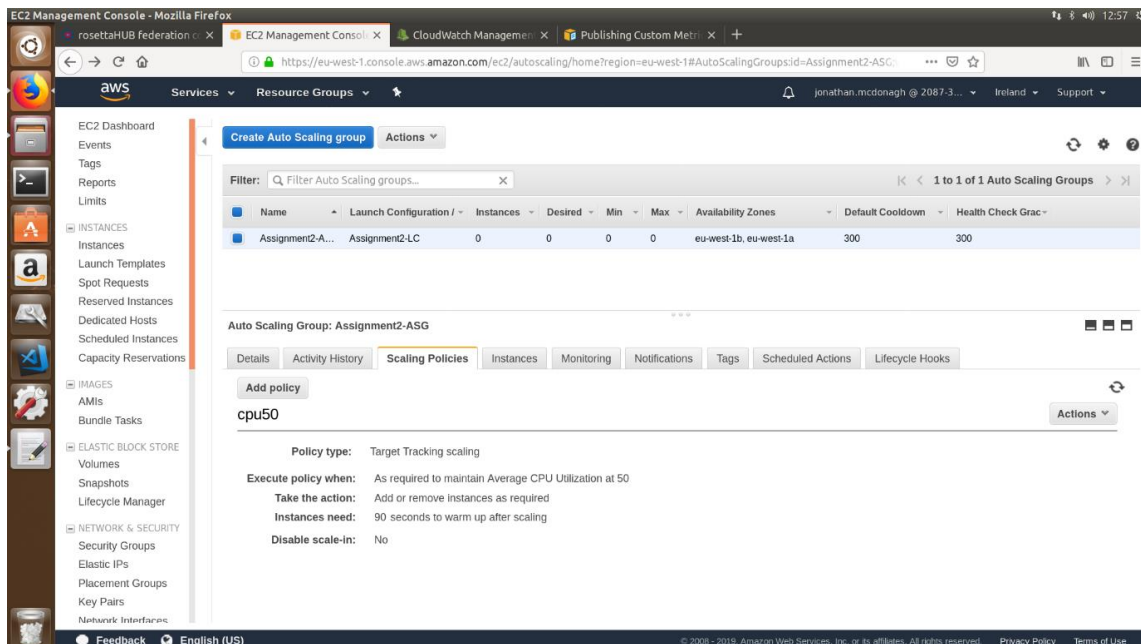


Figure 12: Auto Scaling Policy (cpu50)

## ASSIGNMENT 2

### Step 8 (CloudWatch)

For CloudWatch I used a CPU-Utilization to Increase Group Size when the CPU-Utilization was above 50. I also added a Decrease Group Size when the CPU-Utilization was below 50. I also set the Target Group to my Target Group for my Load Balancer.

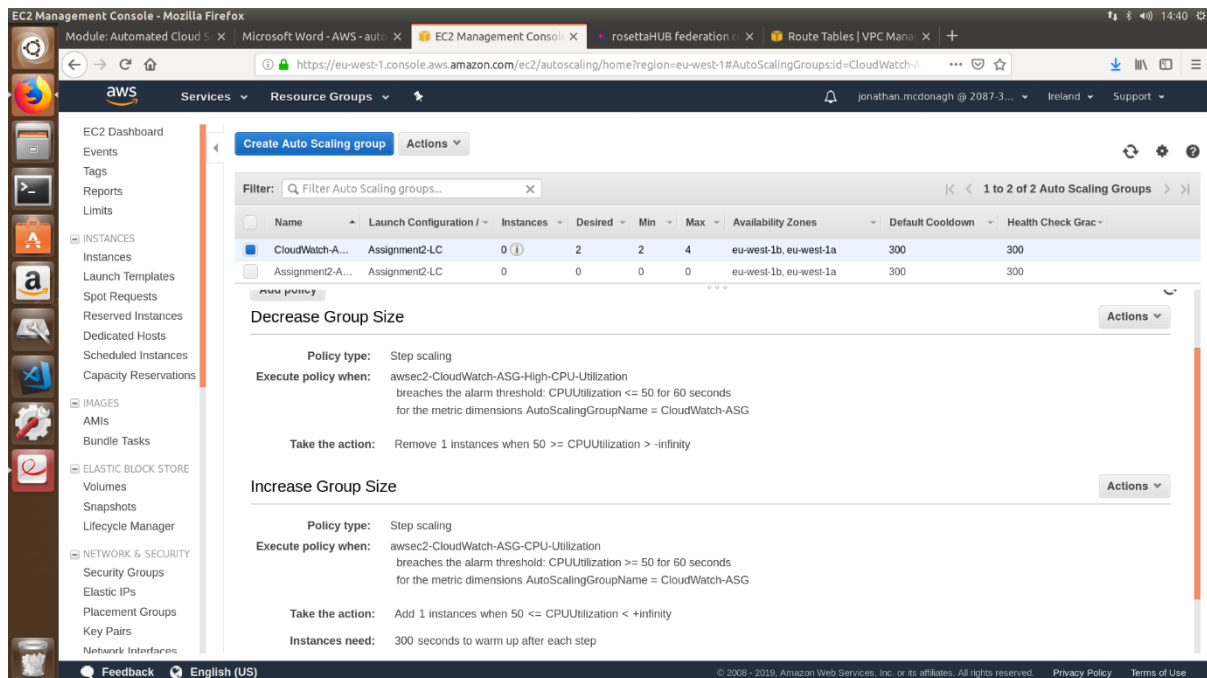


Figure 13: Cloud-Watch Policies

I ran a simple script on one of the Cloud Watch Instances that was created from the Cloud Watch Auto-Scaling Group

*While true; do x=0; done*

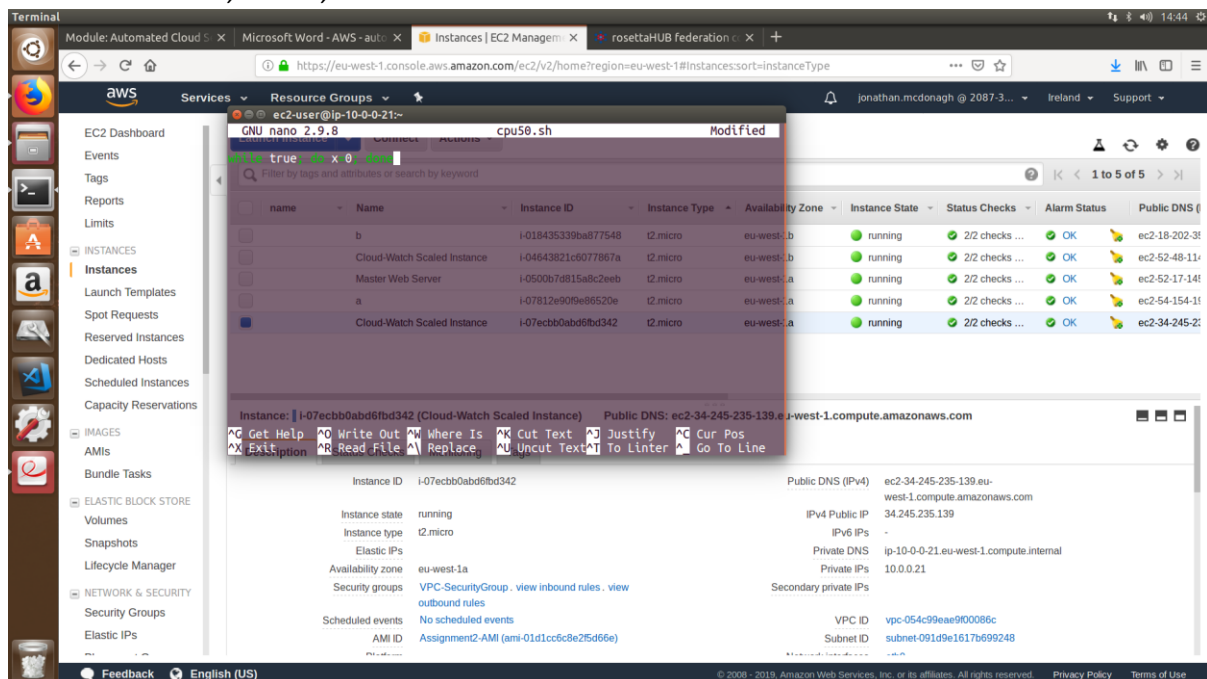


Figure 14: cpu100 script for Cloud Watch

## ASSIGNMENT 2

Here as you can see three instances were created a few minutes after I ran the `cpu100.sh` script.

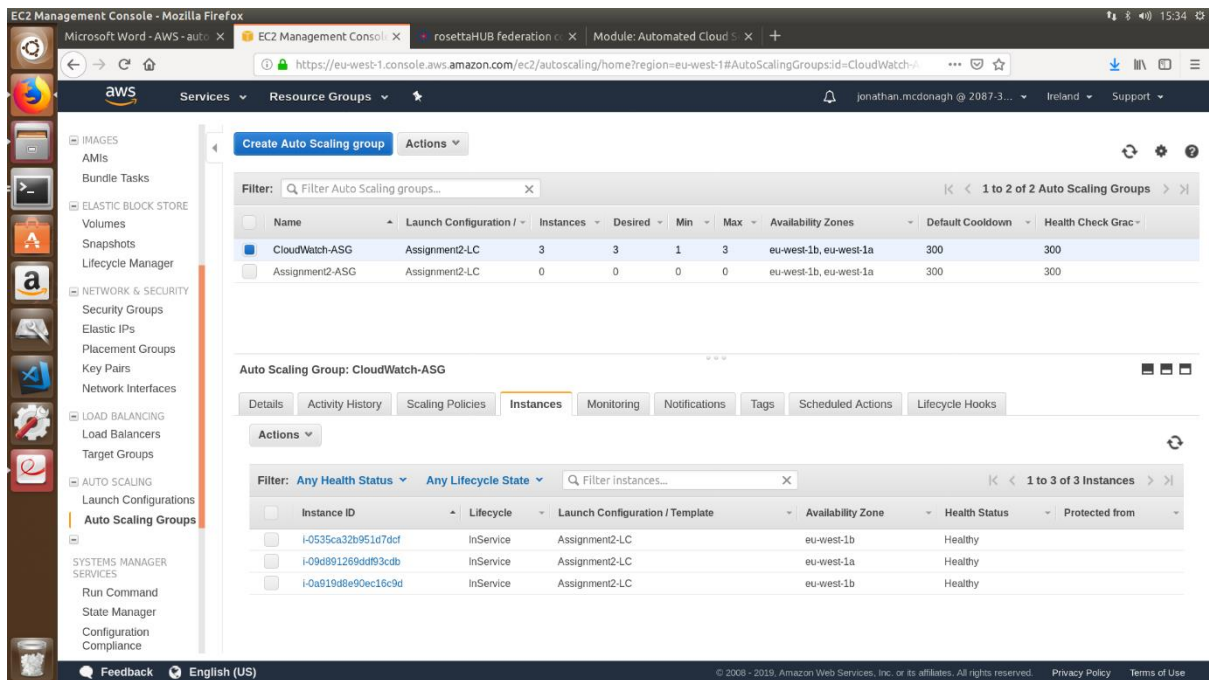


Figure 15: Cloud Watch Auto Scaling

After killing the PID for the `cpu100.sh` script the instances starting terminating until it got back to the min which was 1.

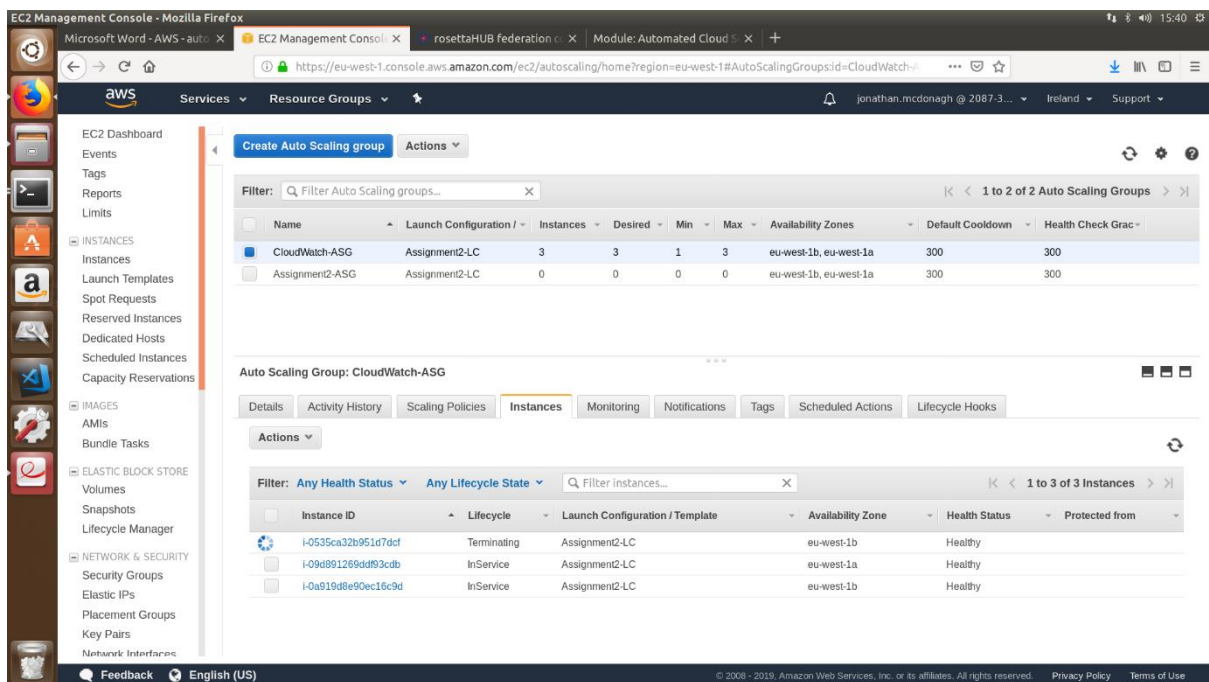


Figure 16: Cloud Watch Instances Terminating



## ASSIGNMENT 2

As you can see below.

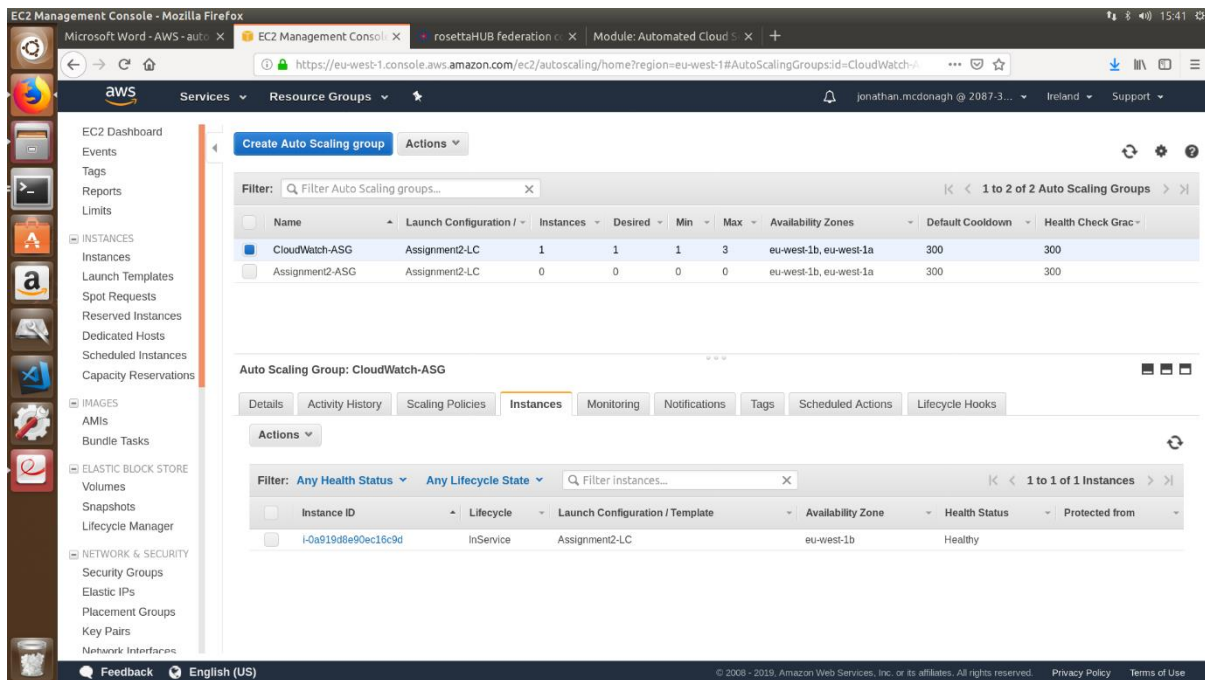


Figure 17: Cloud Watch Back to Min

### Step 9 (Test Traffic to the Load Balancer)

As you can see here I used curl to generate test traffic to my load balancer in order to test my load balancer. The example below sent 100 requests to my load balancer.

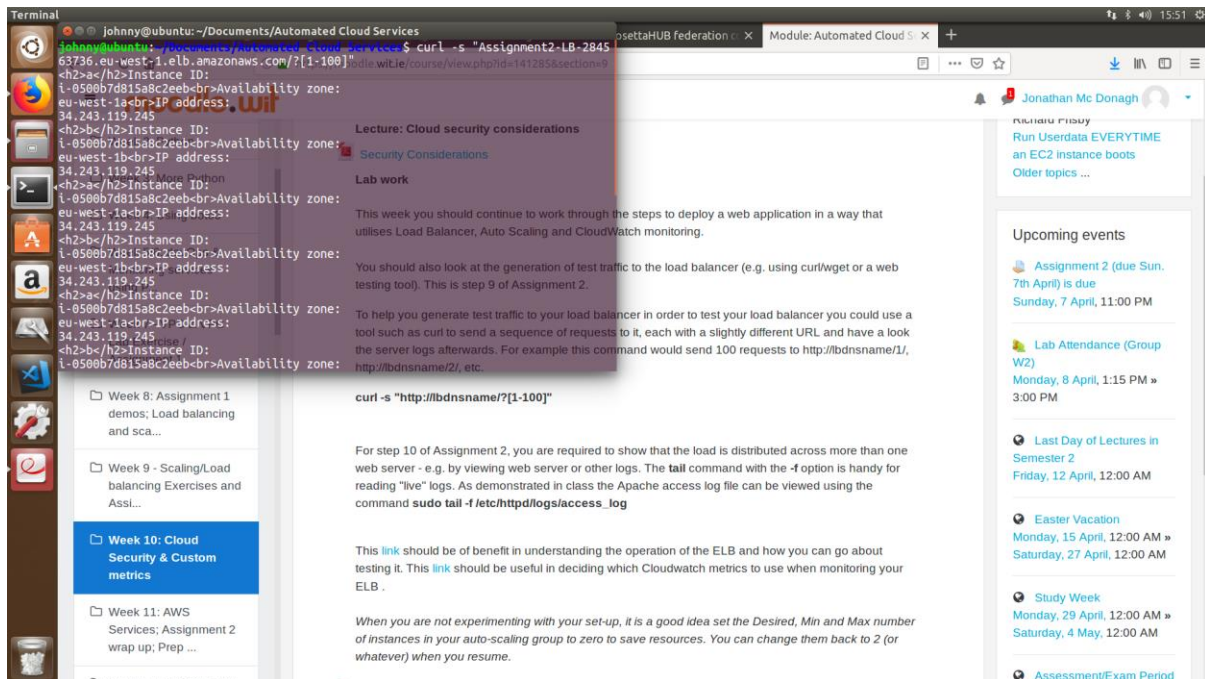


Figure 18: curl command

## Step 10 (To Show the Load is Distributed Across the Web Servers)

As you can see below. I used the curl command to generate test traffic to the Load Balancer. But before this I used the 'sudo tail -f /etc/httpd/logs/access\_log' to view the live logs. As you can see below on the right the numbers are odd and on the left the numbers are even. This shows that the Load Distribution is spread between both of the web servers.

The image shows three terminal windows. The top two windows display Apache access logs from two different EC2 instances. The left window shows logs for IP 10.0.0.34, where the request numbers are even (780, 782, 784, 786, 788, 790, 792, 794, 796, 798, 1000). The right window shows logs for IP 10.0.0.34, where the request numbers are odd (779, 781, 783, 785, 787, 789, 791, 793, 795, 797, 799). The bottom window shows a terminal on a host named 'johnny@ubuntu' where a curl command is used to generate traffic to a load balancer: 'curl -s "Assignment2-LB-284563736.eu-west-1.elb.amazonaws.com/[1-100]"'. Below the command, the output shows the instance ID and IP address for each of the 100 requests, demonstrating that the traffic is distributed across both instances.

```

ec2-user@ip-10-0-0-110:~$ sudo tail -f /etc/httpd/logs/access_log
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /780 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /782 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /784 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /786 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /788 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /790 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /792 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /794 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /796 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /798 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /1000 HTTP/1.1" 200 109 "-" "curl/7.47.0"

ec2-user@ip-10-0-2-171:~$ sudo tail -f /etc/httpd/logs/access_log
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /779 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /781 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /783 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /785 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /787 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /789 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /791 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /793 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /795 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /797 HTTP/1.1" 200 109 "-" "curl/7.47.0"
10.0.0.34 - - [05/Apr/2019:15:51:05 +0000] "GET /799 HTTP/1.1" 200 109 "-" "curl/7.47.0"

johnny@ubuntu:~/Documents/Automated Cloud Services$ curl -s "Assignment2-LB-284563736.eu-west-1.elb.amazonaws.com/[1-100]"
<h2>b</h2>Instance ID:
i-0500b7d815a8c2eeb<br>Availability zone:
eu-west-1a<br>IP address:
34.243.119.245
<h2>a</h2>Instance ID:
i-0500b7d815a8c2eeb<br>Availability zone:
eu-west-1a<br>IP address:
34.243.119.245
<h2>b</h2>Instance ID:
i-0500b7d815a8c2eeb<br>Availability zone:
eu-west-1b<br>IP address:
34.243.119.245
<h2>a</h2>Instance ID:
i-0500b7d815a8c2eeb<br>Availability zone:
eu-west-1a<br>IP address:
34.243.119.245

```

Figure 19: Load Distribution

## Step 11 (Custom CloudWatch Metric)

To monitor some activity on my server I used the put-metric-data API call to push values to CloudWatch. Here is the monitoring script I got from AWS:

```
=====Sample script=====
#!/bin/bash
USEDMEMORY=$(free -m | awk 'NR==2{printf "%.2f\t", $3*100/$2 }')
TCP_CONN=$(netstat -an | wc -l)
TCP_CONN_PORT_80=$(netstat -an | grep 80 | wc -l)
USERS=$(uptime | awk '{ print $6 }')
IO_WAIT=$(iostat | awk 'NR==4 {print $5}')

aws cloudwatch put-metric-data --metric-name memory-usage --dimensions Instance=i-0c51f9f1213e63159 --namespace "Custom" --value $USEDMEMORY
aws cloudwatch put-metric-data --metric-name Tcp_connections --dimensions Instance=i-0c51f9f1213e63159 --namespace "Custom" --value $TCP_CONN
aws cloudwatch put-metric-data --metric-name TCP_connection_on_port_80 --dimensions Instance=i-0c51f9f1213e63159 --namespace "Custom" --value $TCP_CONN_PORT_80
aws cloudwatch put-metric-data --metric-name No_of_users --dimensions Instance=i-0c51f9f1213e63159 --namespace "Custom" --value $USERS
aws cloudwatch put-metric-data --metric-name IO_WAIT --dimensions Instance=i-0c51f9f1213e63159 --namespace "Custom" --value $IO_WAIT
=====
```

Figure 20: AWS Metric Script

To do this I used the Sample Script provided by AWS and saved it as mem.sh and saved it onto my CloudWatch Instance. I then had to use 'aws configure' in the terminal to setup my AWS account's Access Key ID and Secret Access Key. I then copied the mem.sh to the CloudWatch Instance. After this I gave permissions to the file using 'chmod +x mem.sh'.

After this I used 'env EDITOR=nano crontab -e' to create a cron job. I added this line to execute my script every minute, '\*/\*/\*/\*/\* /home/ec2-user/mem.sh'. After this it started installing crontab.

After this I logged into my CloudWatch Metrics and choose Metrics and Custom appeared. I then copied my CloudWatch Instance ID to display the metrics for that Instance. Here is a Screenshot below:

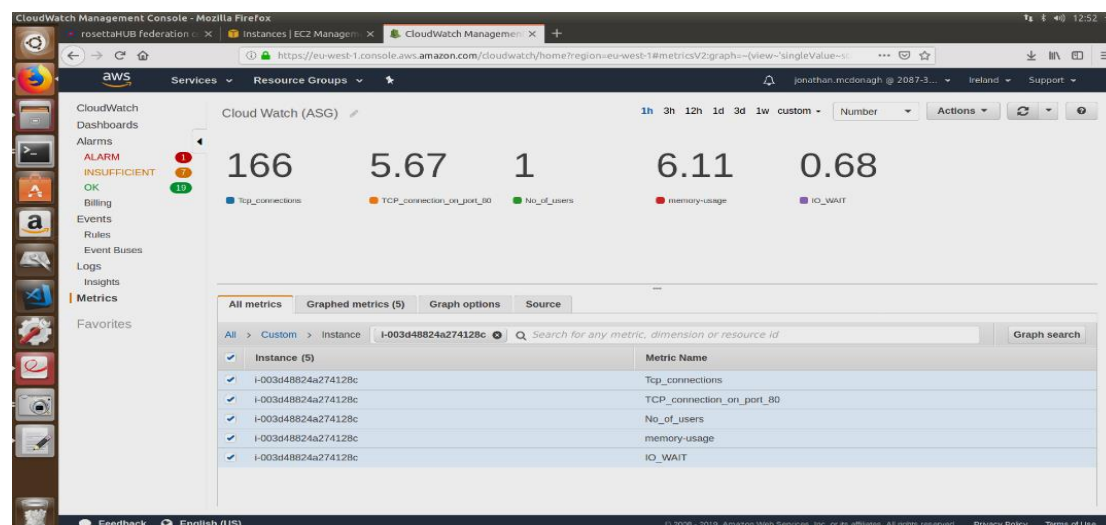


Figure 21: Cloud Watch Monitoring (1)

## ASSIGNMENT 2

After viewing this I launched the `cpu100.sh` script and as you can see below the usage started to increase.

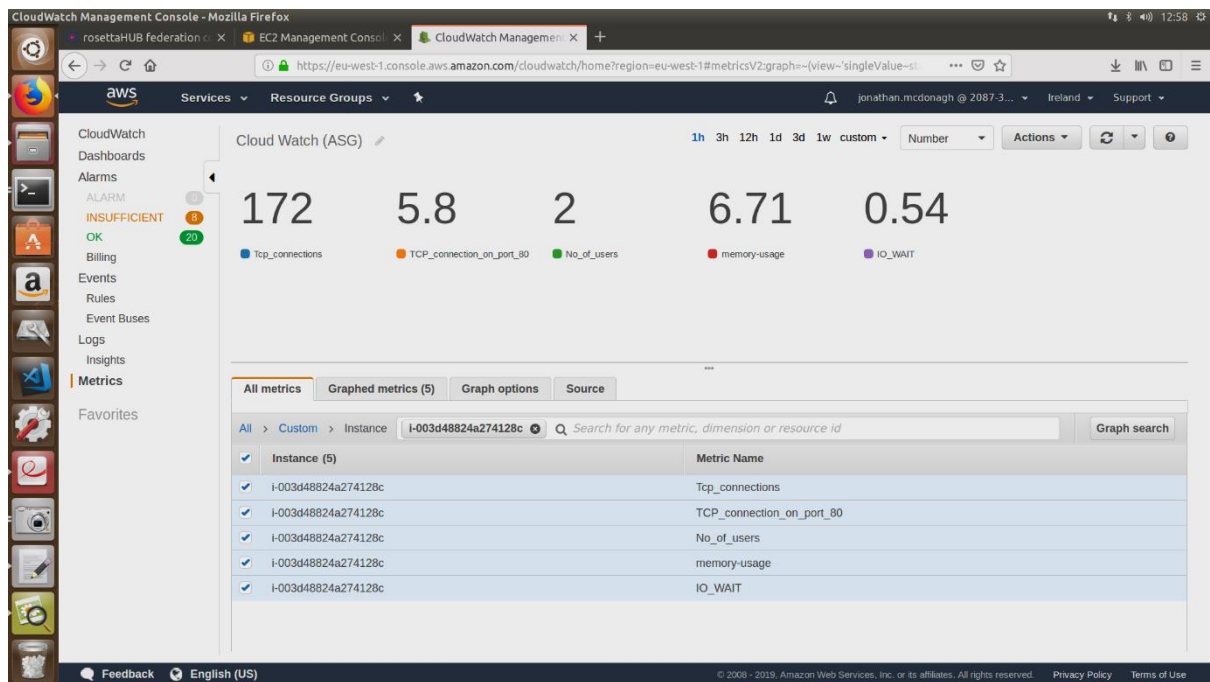


Figure 22: Cloud Watch Monitoring Increased (2)



## References

AWS CloudWatch-Custom-Metrics Metrics. 2019. AWS CloudWatch-Custom-Metrics. [Online]  
Available at: <https://aws.amazon.com/premiumsupport/knowledge-center/cloudwatch-custom-metrics/> [Accessed 04 April 2019].